

Sebastian Dziadzio

- AI researcher specializing in continual pretraining and benchmarking agentic systems
- Ten years of machine learning research experience across startups, industry labs, and academia
- Published at NeurIPS, CVPR, ACL, and ICCV, with 750+ citations, and two patents

 sebastiandziadzio.com |  sebastian.dziadzio@gmail.com |  linkedin.com/in/sebastiandziadzio

EXPERIENCE

Feb 2022–Present	Doctoral Researcher, Bethge Lab, University of Tübingen	Tübingen, Germany
	I developed methods and benchmarks for continually updating multimodal foundation models and evaluating them on open-ended capabilities. More recently, I built agentic benchmarks for inferring mechanisms from behavior.	
Oct 2019–Dec 2021	Scientist, Microsoft Research, Mixed Reality & AI Lab	Cambridge, UK
	I owned and extended a parametric 3D face model used across multiple projects. I built the training, evaluation, and visualization infrastructure, and led a research effort that reduced the average fitting error by half.	
Sep 2018–Sep 2019	AI Resident, Microsoft Research	Cambridge, UK
	I conducted research on GANs for novel view synthesis and on deep RL for video games.	
Sep 2016–Aug 2018	Data Scientist, Cliqz	Munich, Germany
	I built an information retrieval model combining LSTMs and CNNs.	
Mar 2015–Jun 2016	Software Engineer, Nokia Networks	Kraków, Poland
	I implemented LTE control plane features and designed the test suites.	

EDUCATION

Feb 2022–Dec 2026	PhD, Computer Science, University of Tübingen	Tübingen, Germany
	ELLIS PhD Program (3% acceptance rate), IMPRS-IS (5%)	
Mar 2014–Sep 2016	MSc, Computer Science, AGH University	Kraków, Poland
Oct 2010–Feb 2014	BSc, Acoustical Engineering, AGH University	Kraków, Poland

SKILLS

Research areas

Benchmarking and evaluation, continual multimodal pretraining, model merging, agentic systems

Technical skills

Python, PyTorch, Docker, SLURM, DeepSpeed, vLLM

Languages

Polish (native), English (fluent), German (intermediate), Spanish (conversational)

ACADEMIC LEADERSHIP

Organization

ML in PL Conference 2022–2026 (organizer and team lead), ML Summer School 2022–2026

Reviewing

NeurIPS, ICML, ICLR, CVPR

Supervision

One MSc and one PhD student

PUBLICATIONS

* denotes equal first contribution

RevengeBench: Reverse Engineering Code-Space Policies from Behavioral Experiments

NeurIPS 2026, under review

B. Rahmani*, S. Dziadzio*, J. Strüber*, S. Hernández-Gutiérrez, M. Bethge.

ONEBench to Test Them All: Sample-Level Benchmarking Over Open-Ended Capabilities

ACL 2025

A. Ghosh*, S. Dziadzio*, A. Prabhu, V. Udandarao, S. Albanie, M. Bethge.

How to Merge Your Multimodal Models Over Time?

CVPR 2025

S. Dziadzio*, V. Udandarao*, K. Roth*, A. Prabhu, Z. Akata, S. Albanie, M. Bethge.

A Practitioner's Guide to Continual Multimodal Pretraining

NeurIPS 2024

K. Roth*, V. Udandarao*, S. Dziadzio, A. Prabhu, M. Cherti, O. Vinyals, O. Hénaff, S. Albanie, M. Bethge, Z. Akata.

Disentangled Continual Learning: Separating Memory Edits from Generalization

CoLLAs 2024

S. Dziadzio, Ç. Yıldız, G. v/d Ven, T. Trzciński, T. Tuytelaars, M. Bethge.

Full-Body Motion from a Single Head-Mounted Device

ICCV 2021

A. Dittadi, S. Dziadzio, D. Cosker, B. Lundell, T. Cashman, J. Shotton.

Fake It Till You Make It: Face Analysis in the Wild Using Synthetic Data Alone

ICCV 2021

E. Wood, T. Baltrušaitis, C. Hewitt, S. Dziadzio, T. Cashman, J. Shotton.

PATENTS

Controllable Image Generation

US-11748932, 2023

M. Kowalski, S. Garbin, M. Johnson, T. Baltrušaitis, M. De La Gorce, V. Estellers, S. Dziadzio.

Computing Photorealistic Versions of Synthetic Images

US-11354846, 2022

S. Garbin, M. Kowalski, M. Johnson, T. Baltrušaitis, M. De La Gorce, V. Estellers, S. Dziadzio, J. Shotton.